

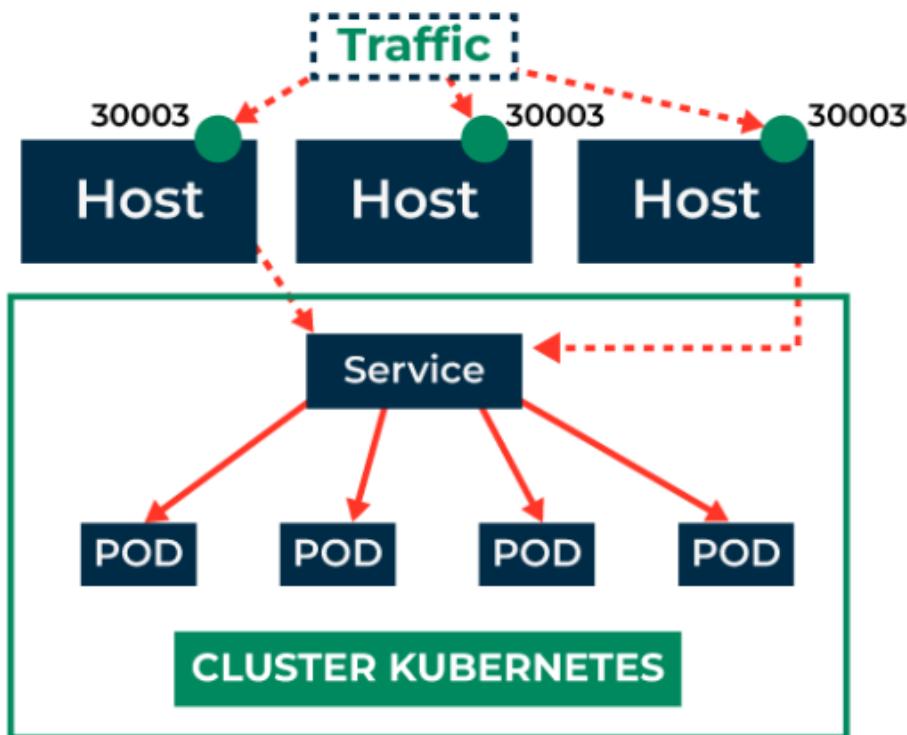
# Kubernetes Services

## Service

Mainly we have 3 types of services.

1. **NodePort**
2. **ClusterIP**
3. **LoadBalancer**

### 1. NodePort



Let's create a service

```
vim service.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: mysvc
spec:
  selector:
    app: webserver
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30080
```

```
kubectl create -f service.yaml
```

```
root@master-node:~# vim service.yaml
root@master-node:~# kubectl create -f service.yaml
service/mysvc created
root@master-node:~#
root@master-node:~# kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP      10d
mysvc      NodePort  10.100.161.39  <none>        80:30080/TCP  17s
root@master-node:~# |
```

```
kubectl describe svc mysvc
```

```
root@master-node:~# kubectl describe svc mysvc
Name:           mysvc
Namespace:      default
Labels:         <none>
Annotations:   <none>
Selector:       app=webserver
Type:          NodePort
IP Family Policy: SingleStack
IP Families:    IPv4
IP:             10.100.161.39
IPs:            10.100.161.39
Port:           <unset>  80/TCP
TargetPort:     80/TCP
NodePort:       <unset>  30080/TCP
Endpoints:     None
Session Affinity:  Cluster
External Traffic Policy: Cluster
Events:        <none>
```

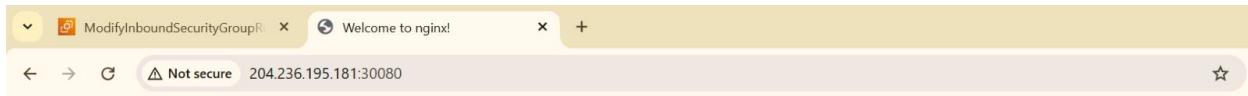
```
kubectl create -f deployment.yaml
```

```
root@master-node:~# kubectl create -f deployment.yaml
deployment.apps/mywebsite created
root@master-node:~#
root@master-node:~#
root@master-node:~# kubectl get deployments.apps
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite  50/50    50          50          29s
root@master-node:~# |
```

Now access our application using worker node public ip and 30080 port

Go browser and access

<http://worker-node-ip/30080>



## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](http://nginx.org).  
Commercial support is available at [nginx.com](http://nginx.com).

*Thank you for using nginx.*

Ok now let's update image in deployment file and replace deployment with our custom image.

vim deployment.yaml

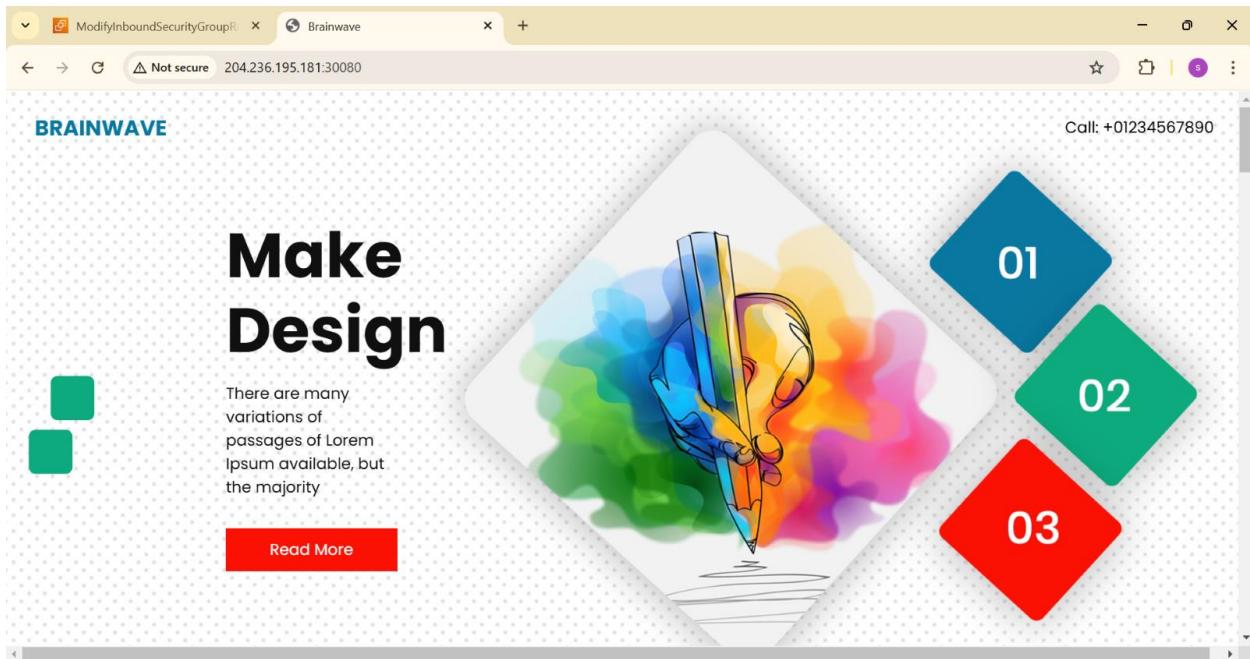
```
kind: Deployment
metadata:
  name: mywebsite
spec:
  selector:
    matchLabels:
      app: webserver
  replicas: 50
  template:
    metadata:
      name: xyz
      labels:
        app: webserver
  spec:
    containers:
      - name: sdfdf
        image: jacksneel/centos-web
```

kubectl replace -f deployment.yaml

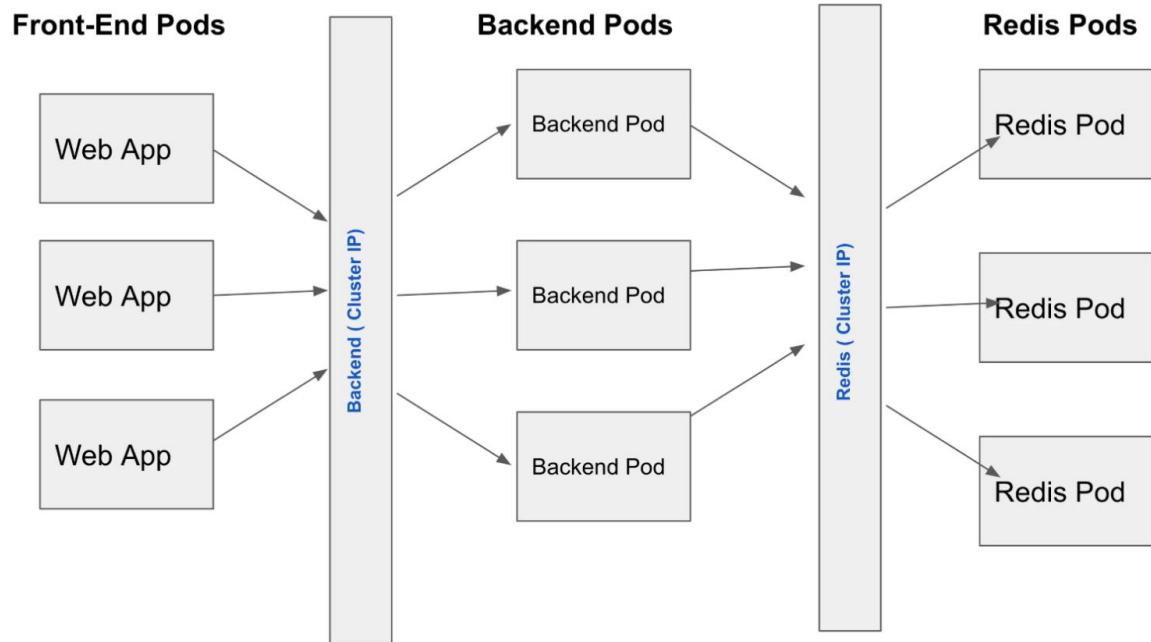
```
root@master-node:~# vim deployment.yaml
root@master-node:~# kubectl replace -f deployment.yaml
deployment.apps/mywebsite replaced
root@master-node:~#
root@master-node:~# kubectl get deployments.apps
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite  38/50   37          38          11m
root@master-node:~#
```

Let's access application

<http://worker-node-ip:30080>



## 2. ClusterIP



Lets create a ClusterIP service for exiting pods

`vim cluster.yml`

```

apiVersion: v1
kind: Service
metadata:
  name: myip
spec:
  selector:
    app: webserver
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80

```

```
kubectl create -f cluster.yml
```

```
root@master-node:~# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP      10d
myip       ClusterIP  10.97.33.109    <none>        80/TCP       7s
mysvc      NodePort   10.100.161.39    <none>        80:30080/TCP  30m
root@master-node:~# |
```

```
root@master-node:~#
root@master-node:~# curl 10.97.33.109
<!DOCTYPE html>
<html lang="en">

<head>
  <!-- Basic -->
  <meta charset="utf-8" />
  <meta http-equiv="X-UA-Compatible" content="IE=edge" />
  <!-- Mobile Metas -->
  <meta name="viewport" content="width=device-width, initial-scale=1, shrink-to-fit=no" />
  <!-- Site Metas -->
```

### 3. LoadBalancer

Exposes the Service externally using an external load balancer. Kubernetes does not directly offer a load balancing component; you must provide one, or you can integrate your Kubernetes cluster with a cloud provider.

type: LoadBalancer

On cloud providers which support external load balancers, setting the type field to LoadBalancer provisions a load balancer for your Service. The actual creation of the load balancer happens asynchronously, and information about the provisioned balancer is published in the Service's .status.loadBalancer field. For example:

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
```

```
app.kubernetes.io/name: MyApp
ports:
- protocol: TCP
  port: 80
  targetPort: 9376
clusterIP: 10.0.171.239
type: LoadBalancer
status:
loadBalancer:
  ingress:
- ip: 192.0.2.127
```

Traffic from the external load balancer is directed at the backend Pods. The cloud provider decides how it is load balanced.

To implement a Service of type: LoadBalancer, Kubernetes typically starts off by making the changes that are equivalent to you requesting a Service of type: NodePort. The cloud-controller-manager component then configures the external load balancer to forward traffic to that assigned node port.

You can configure a load balanced Service to [omit](#) assigning a node port, provided that the cloud provider implementation supports this.

Some cloud providers allow you to specify the `loadBalancerIP`. In those cases, the load-balancer is created with the user-specified `loadBalancerIP`. If the `loadBalancerIP` field is not specified, the load balancer is set up with an ephemeral IP address. If you specify a `loadBalancerIP` but your cloud provider does not support the feature, the `loadbalancerIP` field that you set is ignored.

Traffic from the external load balancer is directed at the backend Pods. The cloud provider decides how it is load balanced.

To implement a Service of type: LoadBalancer, Kubernetes typically starts off by making the changes that are equivalent to you requesting a Service of type: NodePort. The cloud-controller-manager component then configures the external load balancer to forward traffic to that assigned node port.

You can configure a load balanced Service to [omit](#) assigning a node port, provided that the cloud provider implementation supports this.

Some cloud providers allow you to specify the loadBalancerIP. In those cases, the load-balancer is created with the user-specified loadBalancerIP. If the loadBalancerIP field is not specified, the load balancer is set up with an ephemeral IP address. If you specify a loadBalancerIP but your cloud provider does not support the feature, the loadBalancerIP field that you set is ignored.

## Imperative and Declarative Approach in Kubernetes

### Difference Between Imperative and Declarative Approach in Kubernetes

#### Imperative Approach (Command-Based)

##### Definition:

In the **imperative approach**, we **give direct commands** to Kubernetes to perform actions like creating or modifying resources using kubectl run, kubectl create, or kubectl expose.

##### Example:

If we want to create an **Nginx pod** using the imperative approach, we run:

```
kubectl run nginx --image=nginx --port=80
```

##### Or to expose a service:

```
kubectl expose pod nginx --port=80 --type=NodePort
```

##### ◊ Features of Imperative Approach:

- **Direct Commands:** Resources are created and modified using direct commands.
- **Instant Execution:** The commands execute immediately.
- **Less Maintainable:** Changes must be manually applied each time.

## Declarative Approach (YAML-based)

### Definition:

In the **declarative approach**, we **define the desired state** of Kubernetes resources in **YAML files**, and then apply them using kubectl apply. Kubernetes ensures that the system reaches and maintains this desired state.

### Example:

If we want to deploy **Nginx using a Deployment**, we define the following **YAML file**:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```

To apply this YAML file:

```
kubectl apply -f nginx-deployment.yaml
```

### ◊ Features of Declarative Approach:

- **State-Based Management:** We define the **desired state**, and Kubernetes automatically maintains it.
- **Version Control Friendly:** YAML files can be stored in **Git** for tracking changes.
- **Easier Maintenance:** It allows for automation and better management of the cluster.